

Chitramoy_Mukherjee-DSC-630-Project-Final_output 1

June 1, 2024

DSC-630-T302 Chitramoy Mukherjee Date : 05/28/2024 Week9 - DSC-630 Project Milestone-4

Data Preparation :

Data preparation is the first and most important step of model building. By following these data preparation steps, we can ensure that the dataset is clean, relevant, and properly formatted for training a Hotel Booking Cancellation Prediction Model. This will ultimately lead to better model performance and more accurate predictions. Dataset Basic Information is already provided in the earlier milestone where we have provided total Number of Entries, Columns, Data Types and Missing values. Just as a recap, A majority of the columns, 16 to be precise, are of the object data type and 16 columns are of the int64 data type, representing integer values and 4 columns are of the float64 data type, which typically denotes decimal values.

To prepare the data for a Hotel Booking Cancellation Prediction Model, we would typically perform the following steps:

Data Cleaning: 1. Check for missing values in each column and handle them appropriately (e.g., imputation, deletion). 2. Identify and handle any outliers or anomalies in the data. 3. Remove irrelevant columns that do not contribute to the prediction task.

Feature Selection/Engineering: 1. Select relevant features that are likely to influence hotel booking cancellations (e.g., lead_time, arrival_date, stays_in_weekend_nights, stays_in_week_nights, adults, children, babies, market_segment, deposit_type, etc.). 2. Create new features if necessary (e.g., total_stays = stays_in_weekend_nights + stays_in_week_nights).

Handling Categorical Variables: 1. Encode categorical variables into numerical format using techniques like one-hot encoding or label encoding. 2. Handle categorical variables with many unique categories by grouping or binning them.

Data Splitting: 1. Split the dataset into training and testing sets to evaluate the model's performance on unseen data. 2. Optionally, perform stratified sampling to ensure that the distribution of target labels (canceled/not canceled) is similar in both the training and testing sets.

Data Transformation : 1. Apply transformations such as log transformations or Box-Cox transformations to handle skewed distributions in numerical features.

Data Validation: 1. Check the consistency and integrity of the data after preprocessing to ensure that it is suitable for model training.

During the model building step, I will be building Logistic regression, KNN, Decision Tree Classifier, Random Forest Classifier and XgBoost Classifier and will perform a comparison of the model based on F1-score for class '1' (canceled) and will find the best prediction model.

```
[1]: # importing libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import missingno as msno

import warnings
warnings.filterwarnings('ignore')

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, confusion_matrix, \
    classification_report
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import GradientBoostingClassifier
from xgboost import XGBClassifier
from catboost import CatBoostClassifier
from sklearn.ensemble import ExtraTreesClassifier
from lightgbm import LGBMClassifier
from sklearn.ensemble import VotingClassifier

import folium
from folium.plugins import HeatMap
import plotly.express as px
import sort_dataframeby_monthorweek as sd

plt.style.use('fivethirtyeight')
%matplotlib inline
pd.set_option('display.max_columns', 32)

[3]: # Configure Seaborn plot styles: Set background color and use dark grid
sns.set(rc={'axes.facecolor': 'lavender'}, style='darkgrid')

[2]: # reading data
df = pd.read_csv('C:
    ↪\\Users\\Chitramoy\\Desktop\\MS-DSC\\DSC-630\\Week-12\\hotel_bookings.csv')
df.head()

# Data Types:
```

```
# A majority of the columns, 16 to be precise, are of the object data type
↳ (often representing strings or categorical data).
# 16 columns are of the int64 data type, representing integer values.
# 4 columns are of the float64 data type, which typically denotes decimal
↳ values.
# Missing Values:
# The column children has 4 missing values.
# The column country has 488 missing values.
# The column agent has 16,340 missing values.
# The column company has a significant number of missing values, totaling
↳ 112,593

# Based on the data types and the feature explanations provided earlier, we
↳ identified that 20 columns (hotel, is_canceled, arrival_date_year,
↳ arrival_date_month,
# meal, country, market_segment, distribution_channel, is_repeated_guest,
↳ reserved_room_type, assigned_room_type, deposit_type, agent, company,
↳ customer_type,
# reservation_status, name, email, phone-number and credit_card) are
↳ categorical in terms of their semantics. These features must have string
↳ (object)
# data type to ensure proper analysis and interpretation in subsequent steps.
```

```
[2]:      hotel  is_canceled  lead_time  arrival_date_year  arrival_date_month  \
0  Resort Hotel          0        342             2015             July
1  Resort Hotel          0        737             2015             July
2  Resort Hotel          0         7             2015             July
3  Resort Hotel          0         13             2015             July
4  Resort Hotel          0         14             2015             July

      arrival_date_week_number  arrival_date_day_of_month  \
0                             27                         1
1                             27                         1
2                             27                         1
3                             27                         1
4                             27                         1

      stays_in_weekend_nights  stays_in_week_nights  adults  children  babies  \
0                             0                     0       2        0.0      0
1                             0                     0       2        0.0      0
2                             0                     1       1        0.0      0
3                             0                     1       1        0.0      0
4                             0                     2       2        0.0      0

      meal  country  market_segment  distribution_channel  is_repeated_guest  \
0  BB      PRT      Direct          Direct              0
1  BB      PRT      Direct          Direct              0
```

2	BB	GBR	Direct	Direct	0
3	BB	GBR	Corporate	Corporate	0
4	BB	GBR	Online TA	TA/TO	0

	previous_cancellations	previous_bookings_not_canceled	reserved_room_type	\
0	0	0	C	
1	0	0	C	
2	0	0	A	
3	0	0	A	
4	0	0	A	

	assigned_room_type	booking_changes	deposit_type	agent	company	\
0	C	3	No Deposit	NaN	NaN	
1	C	4	No Deposit	NaN	NaN	
2	C	0	No Deposit	NaN	NaN	
3	A	0	No Deposit	304.0	NaN	
4	A	0	No Deposit	240.0	NaN	

	days_in_waiting_list	customer_type	adr	required_car_parking_spaces	\
0	0	Transient	0.0	0	
1	0	Transient	0.0	0	
2	0	Transient	75.0	0	
3	0	Transient	75.0	0	
4	0	Transient	98.0	0	

	total_of_special_requests	reservation_status	reservation_status_date
0	0	Check-Out	2015-07-01
1	0	Check-Out	2015-07-01
2	0	Check-Out	2015-07-02
3	0	Check-Out	2015-07-02
4	1	Check-Out	2015-07-03

```
[3]: df.describe()
```

```
[3]:
```

	is_canceled	lead_time	arrival_date_year	\
count	119390.000000	119390.000000	119390.000000	
mean	0.370416	104.011416	2016.156554	
std	0.482918	106.863097	0.707476	
min	0.000000	0.000000	2015.000000	
25%	0.000000	18.000000	2016.000000	
50%	0.000000	69.000000	2016.000000	
75%	1.000000	160.000000	2017.000000	
max	1.000000	737.000000	2017.000000	

	arrival_date_week_number	arrival_date_day_of_month	\
count	119390.000000	119390.000000	
mean	27.165173	15.798241	

std	13.605138	8.780829
min	1.000000	1.000000
25%	16.000000	8.000000
50%	28.000000	16.000000
75%	38.000000	23.000000
max	53.000000	31.000000

	stays_in_weekend_nights	stays_in_week_nights	adults \
count	119390.000000	119390.000000	119390.000000
mean	0.927599	2.500302	1.856403
std	0.998613	1.908286	0.579261
min	0.000000	0.000000	0.000000
25%	0.000000	1.000000	2.000000
50%	1.000000	2.000000	2.000000
75%	2.000000	3.000000	2.000000
max	19.000000	50.000000	55.000000

	children	babies	is_repeated_guest \
count	119386.000000	119390.000000	119390.000000
mean	0.103890	0.007949	0.031912
std	0.398561	0.097436	0.175767
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000
max	10.000000	10.000000	1.000000

	previous_cancellations	previous_bookings_not_canceled \
count	119390.000000	119390.000000
mean	0.087118	0.137097
std	0.844336	1.497437
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	26.000000	72.000000

	booking_changes	agent	company	days_in_waiting_list \
count	119390.000000	103050.000000	6797.000000	119390.000000
mean	0.221124	86.693382	189.266735	2.321149
std	0.652306	110.774548	131.655015	17.594721
min	0.000000	1.000000	6.000000	0.000000
25%	0.000000	9.000000	62.000000	0.000000
50%	0.000000	14.000000	179.000000	0.000000
75%	0.000000	229.000000	270.000000	0.000000
max	21.000000	535.000000	543.000000	391.000000

	adr	required_car_parking_spaces	total_of_special_requests
count	119390.000000	119390.000000	119390.000000
mean	101.831122	0.062518	0.571363
std	50.535790	0.245291	0.792798
min	-6.380000	0.000000	0.000000
25%	69.290000	0.000000	0.000000
50%	94.575000	0.000000	0.000000
75%	126.000000	0.000000	1.000000
max	5400.000000	8.000000	5.000000

```
[4]: # Dataset Basic Information
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 119390 entries, 0 to 119389
```

```
Data columns (total 32 columns):
```

#	Column	Non-Null Count	Dtype
0	hotel	119390 non-null	object
1	is_canceled	119390 non-null	int64
2	lead_time	119390 non-null	int64
3	arrival_date_year	119390 non-null	int64
4	arrival_date_month	119390 non-null	object
5	arrival_date_week_number	119390 non-null	int64
6	arrival_date_day_of_month	119390 non-null	int64
7	stays_in_weekend_nights	119390 non-null	int64
8	stays_in_week_nights	119390 non-null	int64
9	adults	119390 non-null	int64
10	children	119386 non-null	float64
11	babies	119390 non-null	int64
12	meal	119390 non-null	object
13	country	118902 non-null	object
14	market_segment	119390 non-null	object
15	distribution_channel	119390 non-null	object
16	is_repeated_guest	119390 non-null	int64
17	previous_cancellations	119390 non-null	int64
18	previous_bookings_not_canceled	119390 non-null	int64
19	reserved_room_type	119390 non-null	object
20	assigned_room_type	119390 non-null	object
21	booking_changes	119390 non-null	int64
22	deposit_type	119390 non-null	object
23	agent	103050 non-null	float64
24	company	6797 non-null	float64
25	days_in_waiting_list	119390 non-null	int64
26	customer_type	119390 non-null	object
27	adr	119390 non-null	float64
28	required_car_parking_spaces	119390 non-null	int64
29	total_of_special_requests	119390 non-null	int64

```

30 reservation_status          119390 non-null object
31 reservation_status_date      119390 non-null object
dtypes: float64(4), int64(16), object(12)
memory usage: 29.1+ MB

```

```

[5]: # checking for null values in the dataset
null = pd.DataFrame({'Null Values' : df.isna().sum(), 'Percentage Null Values' :
    ↪ (df.isna().sum()) / (df.shape[0]) * (100)})
null

```

```

[5]:

```


	Null Values	Percentage Null Values
hotel	0	0.000000
is_canceled	0	0.000000
lead_time	0	0.000000
arrival_date_year	0	0.000000
arrival_date_month	0	0.000000
arrival_date_week_number	0	0.000000
arrival_date_day_of_month	0	0.000000
stays_in_weekend_nights	0	0.000000
stays_in_week_nights	0	0.000000
adults	0	0.000000
children	4	0.003350
babies	0	0.000000
meal	0	0.000000
country	488	0.408744
market_segment	0	0.000000
distribution_channel	0	0.000000
is_repeated_guest	0	0.000000
previous_cancellations	0	0.000000
previous_bookings_not_canceled	0	0.000000
reserved_room_type	0	0.000000
assigned_room_type	0	0.000000
booking_changes	0	0.000000
deposit_type	0	0.000000
agent	16340	13.686238
company	112593	94.306893
days_in_waiting_list	0	0.000000
customer_type	0	0.000000
adr	0	0.000000
required_car_parking_spaces	0	0.000000
total_of_special_requests	0	0.000000
reservation_status	0	0.000000
reservation_status_date	0	0.000000

```

[6]: # filling null values with zero
df.fillna(0, inplace = True)

```

```
[7]: # visualizing null values
msno.bar(df)
plt.show()
```



```
[8]: # adults, babies and children cant be zero at same time, so dropping the rows
      ↳having all these zero at same time
filter = (df.children == 0) & (df.adults == 0) & (df.babies == 0)
df[filter]
```

```
[8]:
```

	hotel	is_canceled	lead_time	arrival_date_year	\
2224	Resort Hotel	0	1	2015	
2409	Resort Hotel	0	0	2015	
3181	Resort Hotel	0	36	2015	
3684	Resort Hotel	0	165	2015	
3708	Resort Hotel	0	165	2015	
...	
115029	City Hotel	0	107	2017	
115091	City Hotel	0	1	2017	
116251	City Hotel	0	44	2017	
116534	City Hotel	0	2	2017	
117087	City Hotel	0	170	2017	

	arrival_date_month	arrival_date_week_number	\
2224	October	41	
2409	October	42	
3181	November	47	
3684	December	53	
3708	December	53	

...
115029	June	26
115091	June	26
116251	July	28
116534	July	28
117087	July	30

	arrival_date_day_of_month	stays_in_weekend_nights	\
2224	6	0	
2409	12	0	
3181	20	1	
3684	30	1	
3708	30	2	

...
115029	27	0
115091	30	0
116251	15	1
116534	15	2
117087	27	0

	stays_in_week_nights	adults	children	babies	meal	country	\
2224	3	0	0.0	0	SC	PRT	
2409	0	0	0.0	0	SC	PRT	
3181	2	0	0.0	0	SC	ESP	
3684	4	0	0.0	0	SC	PRT	
3708	4	0	0.0	0	SC	PRT	

...
115029	3	0	0.0	0	BB	CHE
115091	1	0	0.0	0	SC	PRT
116251	1	0	0.0	0	SC	SWE
116534	5	0	0.0	0	SC	RUS
117087	2	0	0.0	0	BB	BRA

	market_segment	distribution_channel	is_repeated_guest	\
2224	Corporate	Corporate	0	
2409	Corporate	Corporate	0	
3181	Groups	TA/TO	0	
3684	Groups	TA/TO	0	
3708	Groups	TA/TO	0	

...
115029	Online TA	TA/TO	0
115091	Complementary	Direct	0
116251	Online TA	TA/TO	0
116534	Online TA	TA/TO	0
117087	Offline TA/TO	TA/TO	0

previous_cancellations	previous_bookings_not_canceled	\
------------------------	--------------------------------	---

2224	0	0
2409	0	0
3181	0	0
3684	0	0
3708	0	0
...
115029	0	0
115091	0	0
116251	0	0
116534	0	0
117087	0	0

	reserved_room_type	assigned_room_type	booking_changes	deposit_type	\
2224	A	I	1	No Deposit	
2409	A	I	0	No Deposit	
3181	A	C	0	No Deposit	
3684	A	A	1	No Deposit	
3708	A	C	1	No Deposit	
...	
115029	A	A	1	No Deposit	
115091	E	K	0	No Deposit	
116251	A	K	2	No Deposit	
116534	A	K	1	No Deposit	
117087	A	A	0	No Deposit	

	agent	company	days_in_waiting_list	customer_type	adr	\
2224	0.0	174.0	0	Transient-Party	0.00	
2409	0.0	174.0	0	Transient	0.00	
3181	38.0	0.0	0	Transient-Party	0.00	
3684	308.0	0.0	122	Transient-Party	0.00	
3708	308.0	0.0	122	Transient-Party	0.00	
...	
115029	7.0	0.0	0	Transient	100.80	
115091	0.0	0.0	0	Transient	0.00	
116251	425.0	0.0	0	Transient	73.80	
116534	9.0	0.0	0	Transient-Party	22.86	
117087	52.0	0.0	0	Transient	0.00	

	required_car_parking_spaces	total_of_special_requests	\
2224	0	0	
2409	0	0	
3181	0	0	
3684	0	0	
3708	0	0	
...	
115029	0	0	
115091	1	1	

116251	0	0
116534	0	1
117087	0	0

	reservation_status	reservation_status_date
2224	Check-Out	2015-10-06
2409	Check-Out	2015-10-12
3181	Check-Out	2015-11-23
3684	Check-Out	2016-01-04
3708	Check-Out	2016-01-05
...
115029	Check-Out	2017-06-30
115091	Check-Out	2017-07-01
116251	Check-Out	2017-07-17
116534	Check-Out	2017-07-22
117087	Check-Out	2017-07-29

[180 rows x 32 columns]

```
[9]: # filter out the rows identified having children, babies and adults as '0'
df = df[~filter]
df
```

```
[9]:
```

	hotel	is_canceled	lead_time	arrival_date_year	\
0	Resort Hotel	0	342	2015	
1	Resort Hotel	0	737	2015	
2	Resort Hotel	0	7	2015	
3	Resort Hotel	0	13	2015	
4	Resort Hotel	0	14	2015	
...	
119385	City Hotel	0	23	2017	
119386	City Hotel	0	102	2017	
119387	City Hotel	0	34	2017	
119388	City Hotel	0	109	2017	
119389	City Hotel	0	205	2017	

	arrival_date_month	arrival_date_week_number	\
0	July	27	
1	July	27	
2	July	27	
3	July	27	
4	July	27	
...	
119385	August	35	
119386	August	35	
119387	August	35	
119388	August	35	

119389

August

35

	arrival_date_day_of_month	stays_in_weekend_nights	\
0	1	0	
1	1	0	
2	1	0	
3	1	0	
4	1	0	
...	
119385	30	2	
119386	31	2	
119387	31	2	
119388	31	2	
119389	29	2	

	stays_in_week_nights	adults	children	babies	meal	country	\
0	0	2	0.0	0	BB	PRT	
1	0	2	0.0	0	BB	PRT	
2	1	1	0.0	0	BB	GBR	
3	1	1	0.0	0	BB	GBR	
4	2	2	0.0	0	BB	GBR	
...	
119385	5	2	0.0	0	BB	BEL	
119386	5	3	0.0	0	BB	FRA	
119387	5	2	0.0	0	BB	DEU	
119388	5	2	0.0	0	BB	GBR	
119389	7	2	0.0	0	HB	DEU	

	market_segment	distribution_channel	is_repeated_guest	\
0	Direct	Direct	0	
1	Direct	Direct	0	
2	Direct	Direct	0	
3	Corporate	Corporate	0	
4	Online TA	TA/TO	0	
...	
119385	Offline TA/TO	TA/TO	0	
119386	Online TA	TA/TO	0	
119387	Online TA	TA/TO	0	
119388	Online TA	TA/TO	0	
119389	Online TA	TA/TO	0	

	previous_cancellations	previous_bookings_not_canceled	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

...
119385	0	0
119386	0	0
119387	0	0
119388	0	0
119389	0	0

	reserved_room_type	assigned_room_type	booking_changes	deposit_type	\
0	C	C	3	No Deposit	
1	C	C	4	No Deposit	
2	A	C	0	No Deposit	
3	A	A	0	No Deposit	
4	A	A	0	No Deposit	
...	
119385	A	A	0	No Deposit	
119386	E	E	0	No Deposit	
119387	D	D	0	No Deposit	
119388	A	A	0	No Deposit	
119389	A	A	0	No Deposit	

	agent	company	days_in_waiting_list	customer_type	adr	\
0	0.0	0.0	0	Transient	0.00	
1	0.0	0.0	0	Transient	0.00	
2	0.0	0.0	0	Transient	75.00	
3	304.0	0.0	0	Transient	75.00	
4	240.0	0.0	0	Transient	98.00	
...	
119385	394.0	0.0	0	Transient	96.14	
119386	9.0	0.0	0	Transient	225.43	
119387	9.0	0.0	0	Transient	157.71	
119388	89.0	0.0	0	Transient	104.40	
119389	9.0	0.0	0	Transient	151.20	

	required_car_parking_spaces	total_of_special_requests	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	1	
...	
119385	0	0	
119386	0	2	
119387	0	4	
119388	0	0	
119389	0	2	

reservation_status	reservation_status_date
--------------------	-------------------------

0	Check-Out	2015-07-01
1	Check-Out	2015-07-01
2	Check-Out	2015-07-02
3	Check-Out	2015-07-02
4	Check-Out	2015-07-03
...
119385	Check-Out	2017-09-06
119386	Check-Out	2017-09-07
119387	Check-Out	2017-09-07
119388	Check-Out	2017-09-07
119389	Check-Out	2017-09-07

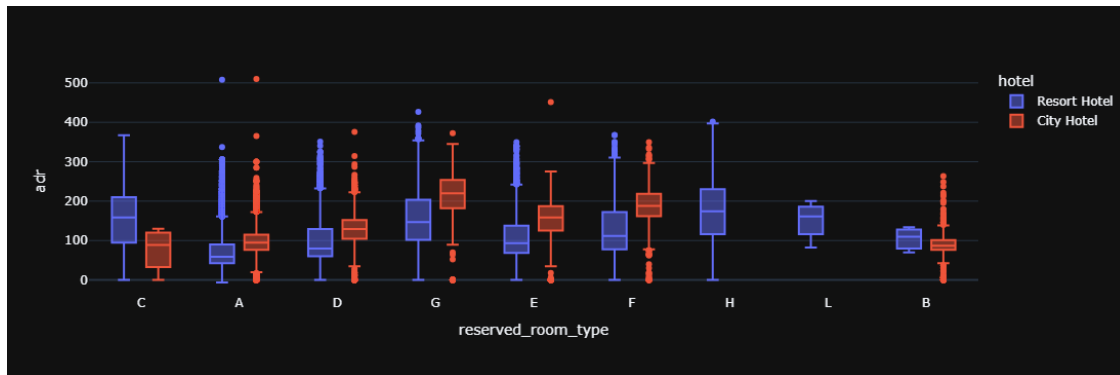
[119210 rows x 32 columns]

```
[15]: # Identify the country wise guests
country_wise_guests = df[df['is_canceled'] == 0]['country'].value_counts().
    ↪reset_index()
country_wise_guests.columns = ['country', 'No of guests']
country_wise_guests
```

```
[15]:      country  No of guests
0      PRT      20977
1      GBR      9668
2      FRA      8468
3      ESP      6383
4      DEU      6067
..      ...      ...
161    BHR          1
162    DJI          1
163    MLI          1
164    NPL          1
165    FRO          1
```

[166 rows x 2 columns]

```
[18]: # Seasonal factors are also important, So the prices varies a lot based on that.
data = df[df['is_canceled'] == 0]
px.box(data_frame = data, x = 'reserved_room_type', y = 'adr', color = 'hotel',
    ↪template = 'plotly_dark')
# The figure shows that the average price per room depends on its type and the
    ↪standard deviation.
```



```
[25]: # we observe here that month column is not in order, and if we visualize we
      ↪ will get improper conclusions.
      # So, first we have to provide right hierarchy to month column
      def sort_month(df, column_name):
          return sd.Sort_Dataframeby_Month(df, column_name)
```

```
[26]: final_prices = sort_month(final_hotel, 'month')
      final_prices
```

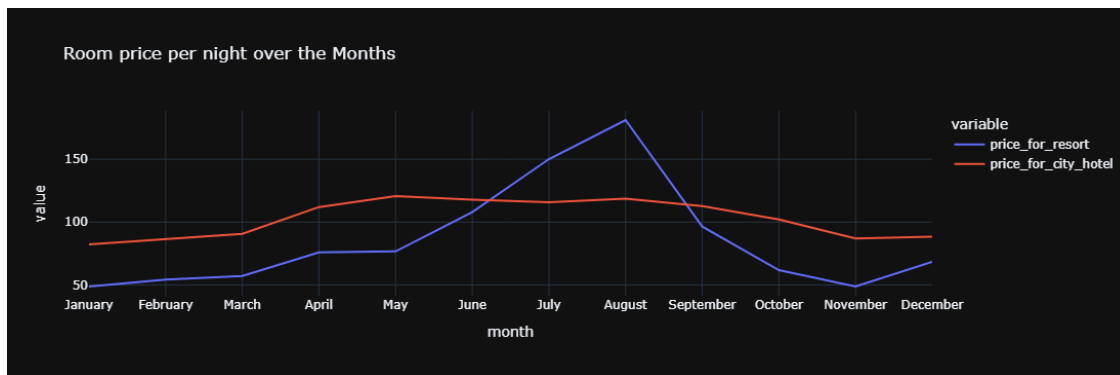
```
[26]:
```

	month	price_for_resort	price_for_city_hotel
0	January	48.761125	82.330983
1	February	54.147478	86.520062
2	March	57.056838	90.658533
3	April	75.867816	111.962267
4	May	76.657558	120.669827
5	June	107.974850	117.874360
6	July	150.122528	115.818019
7	August	181.205892	118.674598
8	September	96.416860	112.776582
9	October	61.775449	102.004672
10	November	48.706289	86.946592
11	December	68.410104	88.401855

```
[27]: plt.figure(figsize = (17, 8))

px.line(final_prices, x = 'month', y =
      ↪ ['price_for_resort', 'price_for_city_hotel'],
      title = 'Room price per night over the Months', template =
      ↪ 'plotly_dark')

# This plot clearly shows that prices in the Resort Hotel are much higher
      ↪ during the summer and prices of city hotel varies less and is most expensive
      ↪ during Spring and Autumn
```



<Figure size 1700x800 with 0 Axes>

```
[40]: # Duration of stay in the hotel.
filter = df['is_canceled'] == 0
data = df[filter]
data.head()
```

```
[40]:
```

	hotel	is_canceled	lead_time	arrival_date_year	arrival_date_month	\
0	Resort Hotel	0	342	2015	July	
1	Resort Hotel	0	737	2015	July	
2	Resort Hotel	0	7	2015	July	
3	Resort Hotel	0	13	2015	July	
4	Resort Hotel	0	14	2015	July	

	arrival_date_week_number	arrival_date_day_of_month	\
0	27	1	
1	27	1	
2	27	1	
3	27	1	
4	27	1	

	stays_in_weekend_nights	stays_in_week_nights	adults	children	babies	\
0	0	0	2	0.0	0	
1	0	0	2	0.0	0	
2	0	1	1	0.0	0	
3	0	1	1	0.0	0	
4	0	2	2	0.0	0	

	meal	country	market_segment	distribution_channel	is_repeated_guest	\
0	BB	PRT	Direct	Direct	0	
1	BB	PRT	Direct	Direct	0	
2	BB	GBR	Direct	Direct	0	
3	BB	GBR	Corporate	Corporate	0	
4	BB	GBR	Online TA	TA/TO	0	

	previous_cancellations	previous_bookings_not_canceled	reserved_room_type	\
0	0	0	C	
1	0	0	C	
2	0	0	A	
3	0	0	A	
4	0	0	A	

	assigned_room_type	booking_changes	deposit_type	agent	company	\
0	C	3	No Deposit	0.0	0.0	
1	C	4	No Deposit	0.0	0.0	
2	C	0	No Deposit	0.0	0.0	
3	A	0	No Deposit	304.0	0.0	
4	A	0	No Deposit	240.0	0.0	

	days_in_waiting_list	customer_type	adr	required_car_parking_spaces	\
0	0	Transient	0.0	0	
1	0	Transient	0.0	0	
2	0	Transient	75.0	0	
3	0	Transient	75.0	0	
4	0	Transient	98.0	0	

	total_of_special_requests	reservation_status	reservation_status_date
0	0	Check-Out	2015-07-01
1	0	Check-Out	2015-07-01
2	0	Check-Out	2015-07-02
3	0	Check-Out	2015-07-02
4	1	Check-Out	2015-07-03

```
[41]: data['total_nights'] = data['stays_in_weekend_nights'] +
      ↪data['stays_in_week_nights']
data.head()
```

```
[41]:      hotel  is_canceled  lead_time  arrival_date_year  arrival_date_month  \
0  Resort Hotel          0        342          2015          July
1  Resort Hotel          0        737          2015          July
2  Resort Hotel          0         7          2015          July
3  Resort Hotel          0        13          2015          July
4  Resort Hotel          0        14          2015          July
```

	arrival_date_week_number	arrival_date_day_of_month	\
0	27	1	
1	27	1	
2	27	1	
3	27	1	
4	27	1	

	stays_in_weekend_nights	stays_in_week_nights	adults	children	babies	\
0	0	0	2	0.0	0	
1	0	0	2	0.0	0	
2	0	1	1	0.0	0	
3	0	1	1	0.0	0	
4	0	2	2	0.0	0	

	meal	country	market_segment	distribution_channel	...	\
0	BB	PRT	Direct	Direct	...	
1	BB	PRT	Direct	Direct	...	
2	BB	GBR	Direct	Direct	...	
3	BB	GBR	Corporate	Corporate	...	
4	BB	GBR	Online TA	TA/TO	...	

	previous_cancellations	previous_bookings_not_canceled	reserved_room_type	\
0	0	0	C	
1	0	0	C	
2	0	0	A	
3	0	0	A	
4	0	0	A	

	assigned_room_type	booking_changes	deposit_type	agent	company	\
0	C	3	No Deposit	0.0	0.0	
1	C	4	No Deposit	0.0	0.0	
2	C	0	No Deposit	0.0	0.0	
3	A	0	No Deposit	304.0	0.0	
4	A	0	No Deposit	240.0	0.0	

	days_in_waiting_list	customer_type	adr	required_car_parking_spaces	\
0	0	Transient	0.0	0	
1	0	Transient	0.0	0	
2	0	Transient	75.0	0	
3	0	Transient	75.0	0	
4	0	Transient	98.0	0	

	total_of_special_requests	reservation_status	reservation_status_date	\
0	0	Check-Out	2015-07-01	
1	0	Check-Out	2015-07-01	
2	0	Check-Out	2015-07-02	
3	0	Check-Out	2015-07-02	
4	1	Check-Out	2015-07-03	

	total_nights
0	0
1	0
2	1
3	1

4

2


[5 rows x 33 columns]

```
[42]: stay = data.groupby(['total_nights', 'hotel']).agg('count').reset_index()
      stay = stay.iloc[:, :3]
      stay = stay.rename(columns={'is_canceled': 'Number of stays'})
      stay
```

```
[42]:   total_nights      hotel  Number of stays
0           0   City Hotel             251
1           0  Resort Hotel             371
2           1   City Hotel            9155
3           1  Resort Hotel            6579
4           2   City Hotel           10983
..          ...          ...
57          46  Resort Hotel              1
58          48   City Hotel              1
59          56  Resort Hotel              1
60          60  Resort Hotel              1
61          69  Resort Hotel              1
```

[62 rows x 3 columns]

```
[44]: px.bar(data_frame = stay, x = 'total_nights', y = 'Number of stays', color = 'hotel',
             barmode = 'group',
             template = 'plotly_dark')
```



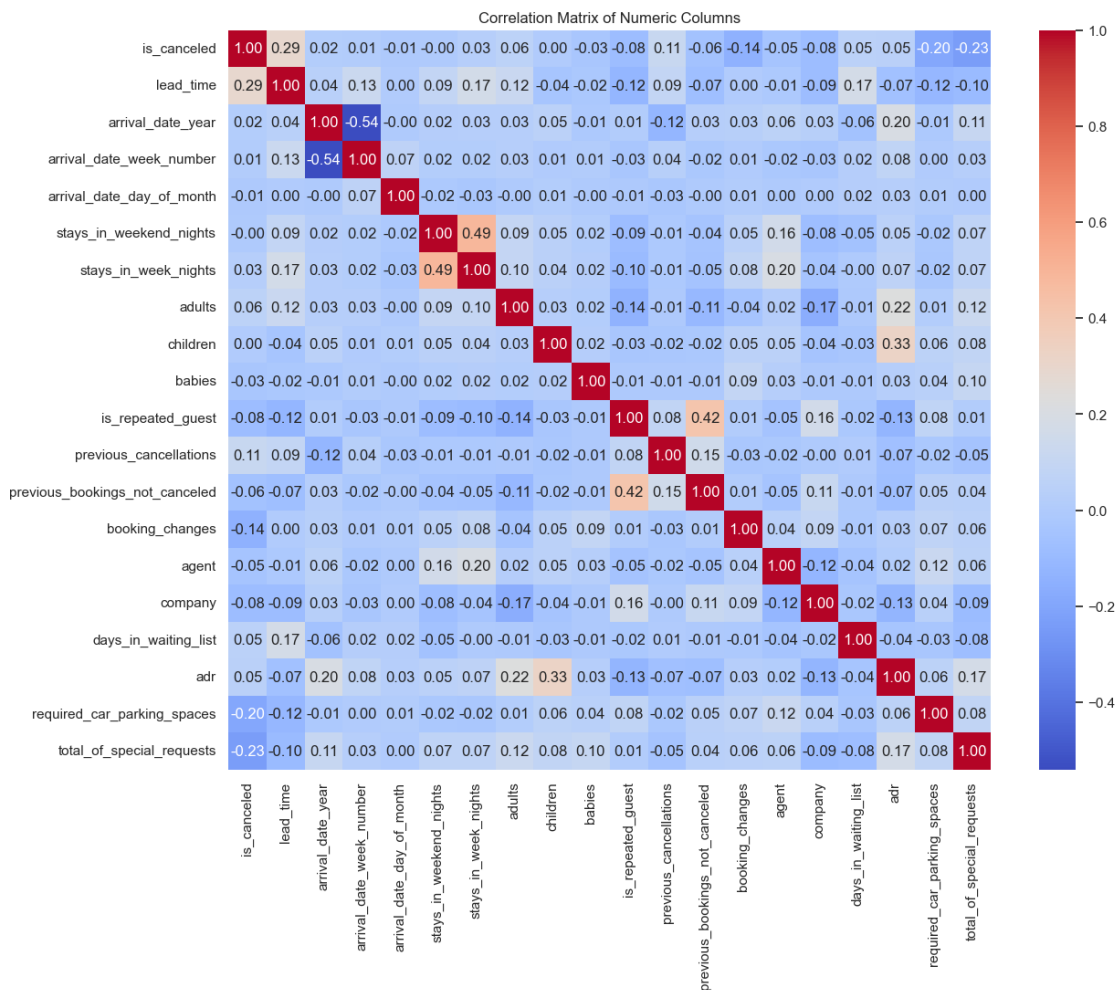
Above plotting shows that for City hotel most of the duration stay is 4 days whereas for Resort Hotel it's 2 days but for Resort hotel 8 days stay is also have significant no. of bookings.

```
[50]: # Identify categorical and numeric columns
      numeric_columns = df.select_dtypes(include=['int64', 'float64']).columns
```

```
# Create correlation matrix for numeric columns
correlation_matrix = df[numeric_columns].corr()
```

```
# Plot the correlation matrix
plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Matrix of Numeric Columns')
plt.show()
```


A correlation matrix of numeric columns provides valuable insights into the relationships between variables in a dataset. By quantifying the strength and direction of linear relationships between pairs of variables, it helps identify potential multicollinearity issues and informs feature selection or dimensionality reduction efforts. Analyzing the correlation matrix enables data scientists to understand how changes in one variable may affect others, facilitating better decision-making in predictive modeling tasks.



```
[16]: # Convert 'arrival_date_month' to categorical type to maintain correct order
months_order = ['January', 'February', 'March', 'April', 'May', 'June', 'July', 'August', 'September', 'October', 'November', 'December']
df['arrival_date_month'] = pd.Categorical(df['arrival_date_month'], categories=months_order, ordered=True)

# Group by 'arrival_date_month' and count bookings and cancellations
monthly_data = df.groupby(['arrival_date_month', 'is_canceled'])['hotel'].count().unstack(fill_value=0)

# Plot
monthly_data.plot(kind='bar', stacked=True, figsize=(12, 6), colormap='coolwarm')
plt.title('Monthly Bookings and Cancellations')
plt.xlabel('Month')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.legend(title='Is Canceled', loc='upper right')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```



From the above plot it shows that month of August have the most no. of bookings as well as having the most no. of cancellations.

```
[15]: # Group by country and count bookings and cancellations
```

```


country_data = df.groupby('country')['is_canceled'].agg([('Total Bookings', 'size'), ('Total Cancellations', 'sum')])

# Sort by total bookings
country_data = country_data.sort_values(by='Total Bookings', ascending=False)

# Select top 5 countries
top_5_countries = country_data.head(5)

# Plot
top_5_countries.plot(kind='bar', figsize=(10, 6))
plt.title('Top 5 Countries by Booking and Cancellation')
plt.xlabel('Country')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.legend(loc='upper right')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()

```



Portugal (PRT) stands out with significantly more reservations than other countries in the top 5, making it the most popular destination among the bookings. It has almost 4 times the reservations to the country in second place, the United Kingdom (GBR). Portugal has the highest cancellation rate among the top 5 countries as shown in the Donut Chart. Almost 56% of its reservations end up being canceled. This indicates that despite Portugal's popularity, it also experiences a relatively high cancellation rate.

```
[53]: # dropping columns that are not useful
```

```
useless_col = ['days_in_waiting_list', 'arrival_date_year',  
               ↪ 'arrival_date_year', 'assigned_room_type', 'booking_changes',  
               'reservation_status', 'country', 'days_in_waiting_list']  
  
df.drop(useless_col, axis = 1, inplace = True)
```

```
[54]: # creating numerical and categorical dataframes
```

```
cat_cols = [col for col in df.columns if df[col].dtype == 'O']  
cat_cols
```

```
[54]: ['hotel',  
       'arrival_date_month',  
       'meal',  
       'market_segment',  
       'distribution_channel',  
       'reserved_room_type',  
       'deposit_type',  
       'customer_type',  
       'reservation_status_date']
```

```
[55]: cat_df = df[cat_cols]  
cat_df.head()
```

```
[55]:      hotel arrival_date_month meal market_segment distribution_channel \  
0  Resort Hotel          July    BB          Direct          Direct  
1  Resort Hotel          July    BB          Direct          Direct  
2  Resort Hotel          July    BB          Direct          Direct  
3  Resort Hotel          July    BB    Corporate    Corporate  
4  Resort Hotel          July    BB    Online TA    TA/TO  
  
   reserved_room_type deposit_type customer_type reservation_status_date  
0                  C   No Deposit    Transient          2015-07-01  
1                  C   No Deposit    Transient          2015-07-01  
2                  A   No Deposit    Transient          2015-07-02  
3                  A   No Deposit    Transient          2015-07-02  
4                  A   No Deposit    Transient          2015-07-03
```

```
[56]: cat_df['reservation_status_date'] = pd.  
       ↪to_datetime(cat_df['reservation_status_date'])  
  
cat_df['year'] = cat_df['reservation_status_date'].dt.year  
cat_df['month'] = cat_df['reservation_status_date'].dt.month  
cat_df['day'] = cat_df['reservation_status_date'].dt.day
```

```
[57]: cat_df.drop(['reservation_status_date', 'arrival_date_month'] , axis = 1,
↳ inplace = True)
```

```
[58]: cat_df.head()
```

```
[58]:      hotel meal market_segment distribution_channel reserved_room_type \
0  Resort Hotel  BB      Direct      Direct      C
1  Resort Hotel  BB      Direct      Direct      C
2  Resort Hotel  BB      Direct      Direct      A
3  Resort Hotel  BB  Corporate  Corporate      A
4  Resort Hotel  BB  Online TA      TA/TO      A

  deposit_type customer_type  year  month  day
0  No Deposit      Transient  2015     7    1
1  No Deposit      Transient  2015     7    1
2  No Deposit      Transient  2015     7    2
3  No Deposit      Transient  2015     7    2
4  No Deposit      Transient  2015     7    3
```

```
[59]: # printing unique values of each column
for col in cat_df.columns:
    print(f"{col}: \n{cat_df[col].unique()}\n")
```

hotel:

```
['Resort Hotel' 'City Hotel']
```

meal:

```
['BB' 'FB' 'HB' 'SC' 'Undefined']
```

market_segment:

```
['Direct' 'Corporate' 'Online TA' 'Offline TA/TO' 'Complementary' 'Groups'
 'Undefined' 'Aviation']
```

distribution_channel:

```
['Direct' 'Corporate' 'TA/TO' 'Undefined' 'GDS']
```

reserved_room_type:

```
['C' 'A' 'D' 'E' 'G' 'F' 'H' 'L' 'B']
```

deposit_type:

```
['No Deposit' 'Refundable' 'Non Refund']
```

customer_type:

```
['Transient' 'Contract' 'Transient-Party' 'Group']
```

year:

```
[2015 2014 2016 2017]
```


month:

```
[ 7  5  4  6  3  8  9  1 11 10 12  2]
```

day:

```
[ 1  2  3  6 22 23  5  7  8 11 15 16 29 19 18  9 13  4 12 26 17 10 20 14
 30 28 25 21 27 24 31]
```

[60]: *# encoding categorical variables*

```
cat_df['hotel'] = cat_df['hotel'].map({'Resort Hotel' : 0, 'City Hotel' : 1})

cat_df['meal'] = cat_df['meal'].map({'BB' : 0, 'FB': 1, 'HB': 2, 'SC': 3,
    ↪ 'Undefined': 4})

cat_df['market_segment'] = cat_df['market_segment'].map({'Direct': 0,
    ↪ 'Corporate': 1, 'Online TA': 2, 'Offline TA/TO': 3,
    ↪ 'Complementary': 4,
    ↪ 'Groups': 5, 'Undefined': 6, 'Aviation': 7})

cat_df['distribution_channel'] = cat_df['distribution_channel'].map({'Direct':
    ↪ 0, 'Corporate': 1, 'TA/TO': 2, 'Undefined': 3,
    ↪ 'GDS':
    ↪ 4})

cat_df['reserved_room_type'] = cat_df['reserved_room_type'].map({'C': 0, 'A':
    ↪ 1, 'D': 2, 'E': 3, 'G': 4, 'F': 5, 'H': 6,
    ↪ 'L': 7, 'B':
    ↪ 8})

cat_df['deposit_type'] = cat_df['deposit_type'].map({'No Deposit': 0,
    ↪ 'Refundable': 1, 'Non Refund': 3})

cat_df['customer_type'] = cat_df['customer_type'].map({'Transient': 0,
    ↪ 'Contract': 1, 'Transient-Party': 2, 'Group': 3})

cat_df['year'] = cat_df['year'].map({2015: 0, 2014: 1, 2016: 2, 2017: 3})
```

[61]: *# Display the categorical variables output*

```
cat_df.head()
```

```
[61]:   hotel  meal  market_segment  distribution_channel  reserved_room_type  \
0      0     0                0                0          0
1      0     0                0                0          0
2      0     0                0                0          1
3      0     0                1                1          1
4      0     0                2                2          1
```

	deposit_type	customer_type	year	month	day
0	0	0	0	7	1
1	0	0	0	7	1
2	0	0	0	7	2
3	0	0	0	7	2
4	0	0	0	7	3

```
[62]: # drop the categorical columns and select and view numerical columns
num_df = df.drop(columns = cat_cols, axis = 1)
num_df.drop('is_canceled', axis = 1, inplace = True)
num_df
```

```
[62]:      lead_time  arrival_date_week_number  arrival_date_day_of_month \
0          342                        27                        1
1          737                        27                        1
2           7                        27                        1
3          13                        27                        1
4          14                        27                        1
...
119385      23                        35                        30
119386     102                        35                        31
119387      34                        35                        31
119388     109                        35                        31
119389     205                        35                        29
```

	stays_in_weekend_nights	stays_in_week_nights	adults	children	\
0	0	0	2	0.0	
1	0	0	2	0.0	
2	0	1	1	0.0	
3	0	1	1	0.0	
4	0	2	2	0.0	
...	
119385	2	5	2	0.0	
119386	2	5	3	0.0	
119387	2	5	2	0.0	
119388	2	5	2	0.0	
119389	2	7	2	0.0	

	babies	is_repeated_guest	previous_cancellations	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	
...	
119385	0	0	0	

119386	0	0	0
119387	0	0	0
119388	0	0	0
119389	0	0	0

	previous_bookings_not_canceled	agent	company	adr	\
0	0	0.0	0.0	0.00	
1	0	0.0	0.0	0.00	
2	0	0.0	0.0	75.00	
3	0	304.0	0.0	75.00	
4	0	240.0	0.0	98.00	
...		
119385	0	394.0	0.0	96.14	
119386	0	9.0	0.0	225.43	
119387	0	9.0	0.0	157.71	
119388	0	89.0	0.0	104.40	
119389	0	9.0	0.0	151.20	

	required_car_parking_spaces	total_of_special_requests
0	0	0
1	0	0
2	0	0
3	0	0
4	0	1
...
119385	0	0
119386	0	2
119387	0	4
119388	0	0
119389	0	2

[119210 rows x 16 columns]

```
[63]: # Display the numerical variables output
num_df.var()
```

```
[63]: lead_time          11422.361808
arrival_date_week_number  184.990111
arrival_date_day_of_month  77.107192
stays_in_weekend_nights   0.990258
stays_in_week_nights      3.599010
adults                    0.330838
children                  0.159070
babies                    0.009508
is_repeated_guest         0.030507
previous_cancellations     0.713887
previous_bookings_not_canceled 2.244415
```

```

agent                11485.169679
company              2897.684308
adr                 2543.589039
required_car_parking_spaces  0.060201
total_of_special_requests  0.628652
dtype: float64

```

```
[64]: # normalizing numerical variables
```

```

num_df['lead_time'] = np.log(num_df['lead_time'] + 1)
num_df['arrival_date_week_number'] = np.log(num_df['arrival_date_week_number'] + 1)
num_df['arrival_date_day_of_month'] = np.
    log(num_df['arrival_date_day_of_month'] + 1)
num_df['agent'] = np.log(num_df['agent'] + 1)
num_df['company'] = np.log(num_df['company'] + 1)
num_df['adr'] = np.log(num_df['adr'] + 1)

```

```
[65]: num_df.var()
```

```

[65]: lead_time                2.582757
arrival_date_week_number      0.440884
arrival_date_day_of_month     0.506325
stays_in_weekend_nights       0.990258
stays_in_week_nights          3.599010
adults                        0.330838
children                      0.159070
babies                        0.009508
is_repeated_guest             0.030507
previous_cancellations        0.713887
previous_bookings_not_canceled 2.244415
agent                         3.535793
company                       1.346883
adr                           0.515480
required_car_parking_spaces    0.060201
total_of_special_requests      0.628652
dtype: float64

```

```
[66]: num_df['adr'] = num_df['adr'].fillna(value = num_df['adr'].mean())
```

```

[67]: X = pd.concat([cat_df, num_df], axis = 1)
      y = df['is_canceled']

```

```
[68]: X.shape, y.shape
```

```
[68]: ((119210, 26), (119210,))
```

```
[69]: # splitting data into training set and test set
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.30)
```

```
[70]: X_train.head()
```

```
[70]:
```

	hotel	meal	market_segment	distribution_channel	reserved_room_type	\
30408	0	0	5	0	1	
84226	1	3	2	2	1	
115183	1	0	2	2	1	
62917	1	0	5	2	1	
54644	1	0	0	0	5	

	deposit_type	customer_type	year	month	day	lead_time	\
30408	0	2	2	11	21	4.753590	
84226	0	2	2	2	21	1.945910	
115183	0	0	3	7	3	5.135798	
62917	3	0	2	11	25	4.127134	
54644	0	0	2	1	18	5.267858	

	arrival_date_week_number	arrival_date_day_of_month	\
30408	3.871201	2.890372	
84226	2.197225	3.044522	
115183	3.332205	1.098612	
62917	1.609438	3.258097	
54644	3.433987	3.135494	

	stays_in_weekend_nights	stays_in_week_nights	adults	children	\
30408	1	3	2	0.0	
84226	0	1	2	0.0	
115183	1	0	1	0.0	
62917	0	2	2	0.0	
54644	0	1	2	2.0	

	babies	is_repeated_guest	previous_cancellations	\
30408	0	0	0	
84226	0	0	0	
115183	0	0	0	
62917	0	0	0	
54644	1	0	0	

	previous_bookings_not_canceled	agent	company	adr	\
30408	0	0.000000	0.0	3.555348	
84226	0	2.302585	0.0	4.394449	
115183	0	2.302585	0.0	4.596129	
62917	0	5.789960	0.0	4.394449	
54644	0	2.708050	0.0	5.152713	

	required_car_parking_spaces	total_of_special_requests
30408	0	0
84226	0	1
115183	0	0
62917	0	0
54644	0	1

```
[71]: X_test.head()
```

```
[71]:
```

	hotel	meal	market_segment	distribution_channel	reserved_room_type	\
41128	1	0	3	2	1	
112779	1	0	2	2	1	
27244	0	0	1	1	1	
98610	1	0	2	2	2	
78196	1	0	0	0	1	

	deposit_type	customer_type	year	month	day	lead_time	\
41128	0	0	0	8	11	0.693147	
112779	0	2	3	5	28	4.859812	
27244	0	0	2	8	24	1.609438	
98610	0	0	2	10	2	4.875197	
78196	0	0	0	10	6	2.564949	

	arrival_date_week_number	arrival_date_day_of_month	\
41128	3.526361	2.564949	
112779	3.091042	3.295837	
27244	3.583519	3.178054	
98610	3.713572	3.401197	
78196	3.713572	1.386294	

	stays_in_weekend_nights	stays_in_week_nights	adults	children	\
41128	0	4	2	0.0	
112779	0	2	2	0.0	
27244	0	1	1	0.0	
98610	0	3	2	0.0	
78196	2	1	2	0.0	

	babies	is_repeated_guest	previous_cancellations	\
41128	0	0	0	
112779	0	0	0	
27244	0	0	0	
98610	0	0	0	
78196	0	0	0	

	previous_bookings_not_canceled	agent	company	adr	\
41128	0	3.761200	0.0000	4.709530	

112779	0	2.302585	0.0000	4.844187
27244	0	0.000000	3.7612	4.934474
98610	0	2.302585	0.0000	4.945207
78196	0	2.708050	0.0000	4.943997

	required_car_parking_spaces	total_of_special_requests
41128	0	0
112779	0	1
27244	1	0
98610	0	2
78196	0	0

```
[72]: y_train.head(), y_test.head()
```

```
[72]: (30408    0
      84226    0
      115183   0
      62917    1
      54644    1
      Name: is_canceled, dtype: int64,
      41128    1
      112779    0
      27244    0
      98610    0
      78196    0
      Name: is_canceled, dtype: int64)
```

```
[73]: # LR model
lr = LogisticRegression()
lr.fit(X_train, y_train)

y_pred_lr = lr.predict(X_test)

acc_lr = accuracy_score(y_test, y_pred_lr)
conf = confusion_matrix(y_test, y_pred_lr)
clf_report = classification_report(y_test, y_pred_lr)

print(f"Accuracy Score of Logistic Regression is : {acc_lr}")
print(f"Confusion Matrix : \n{conf}")
print(f"Classification Report : \n{clf_report}")
```

```
Accuracy Score of Logistic Regression is : 0.8092721527836032
Confusion Matrix :
[[21263  1170]
 [ 5651  7679]]
Classification Report :
              precision    recall  f1-score   support
```

0	0.79	0.95	0.86	22433
1	0.87	0.58	0.69	13330
accuracy				0.81 35763
macro avg	0.83	0.76	0.78	35763
weighted avg	0.82	0.81	0.80	35763

Accuracy Score: 0.8093 Precision for class 0 (not canceled) is 0.79, indicating that 79% of the instances predicted as not canceled were actually not canceled. Recall for class 0 is 0.95, meaning that 95% of the actual not canceled instances were correctly predicted. Precision for class 1 (canceled) is 0.87, indicating that 87% of the instances predicted as canceled were actually canceled. Recall for class 1 is 0.58, meaning that 58% of the actual canceled instances were correctly predicted.

```
[79]: # KNN model
knn = KNeighborsClassifier()
knn.fit(X_train, y_train)

y_pred_knn = knn.predict(X_test)

acc_knn = accuracy_score(y_test, y_pred_knn)
conf = confusion_matrix(y_test, y_pred_knn)
clf_report = classification_report(y_test, y_pred_knn)

print(f"Accuracy Score of KNN is : {acc_knn}")
print(f"Confusion Matrix : \n{conf}")
print(f"Classification Report : \n{clf_report}")
```

Accuracy Score of KNN is : 0.8903056231300506

Confusion Matrix :

```
[[21660  773]
 [ 3150 10180]]
```

Classification Report :

	precision	recall	f1-score	support
0	0.87	0.97	0.92	22433
1	0.93	0.76	0.84	13330
accuracy			0.89	35763
macro avg	0.90	0.86	0.88	35763
weighted avg	0.89	0.89	0.89	35763

Accuracy Score: 0.8903 Precision for class 0 is 0.87, indicating that 87% of the instances predicted as not canceled were actually not canceled. Recall for class 0 is 0.97, meaning that 97% of the actual not canceled instances were correctly predicted. Precision for class 1 is 0.93, indicating that 93% of the instances predicted as canceled were actually canceled. Recall for class 1 is 0.76, meaning that 76% of the actual canceled instances were correctly predicted.


```
[75]: # DT model
dtc = DecisionTreeClassifier()
dtc.fit(X_train, y_train)

y_pred_dtc = dtc.predict(X_test)

acc_dtc = accuracy_score(y_test, y_pred_dtc)
conf = confusion_matrix(y_test, y_pred_dtc)
clf_report = classification_report(y_test, y_pred_dtc)

print(f"Accuracy Score of Decision Tree is : {acc_dtc}")
print(f"Confusion Matrix : \n{conf}")
print(f"Classification Report : \n{clf_report}")
```

Accuracy Score of Decision Tree is : 0.946676732936275

Confusion Matrix :

```
[[21462  971]
 [ 936 12394]]
```

Classification Report :

	precision	recall	f1-score	support
0	0.96	0.96	0.96	22433
1	0.93	0.93	0.93	13330
accuracy			0.95	35763
macro avg	0.94	0.94	0.94	35763
weighted avg	0.95	0.95	0.95	35763

Accuracy Score: 0.9467 Precision for class 0 is 0.96, indicating that 96% of the instances predicted as not canceled were actually not canceled. Recall for class 0 is 0.96, meaning that 96% of the actual not canceled instances were correctly predicted. Precision for class 1 is 0.93, indicating that 93% of the instances predicted as canceled were actually canceled. Recall for class 1 is 0.93, meaning that 93% of the actual canceled instances were correctly predicted.

Additionally, the model is not overfitting, as the metric values for the test and train sets are close together, indicating that the model is generalizing well to unseen data.

```
[76]: # Random Forest model
rd_clf = RandomForestClassifier()
rd_clf.fit(X_train, y_train)

y_pred_rd_clf = rd_clf.predict(X_test)

acc_rd_clf = accuracy_score(y_test, y_pred_rd_clf)
conf = confusion_matrix(y_test, y_pred_rd_clf)
clf_report = classification_report(y_test, y_pred_rd_clf)

print(f"Accuracy Score of Random Forest is : {acc_rd_clf}")
```

```
print(f"Confusion Matrix : \n{conf}")
print(f"Classification Report : \n{clf_report}")
```

Accuracy Score of Random Forest is : 0.9547017867628554

Confusion Matrix :

```
[[22270  163]
 [ 1457 11873]]
```

Classification Report :

	precision	recall	f1-score	support
0	0.94	0.99	0.96	22433
1	0.99	0.89	0.94	13330
accuracy			0.95	35763
macro avg	0.96	0.94	0.95	35763
weighted avg	0.96	0.95	0.95	35763

Accuracy Score: 0.9547 Precision for class 0 is 0.94, indicating that 94% of the instances predicted as not canceled were actually not canceled. Recall for class 0 is 0.99, meaning that 99% of the actual not canceled instances were correctly predicted. Precision for class 1 is 0.99, indicating that 99% of the instances predicted as canceled were actually canceled. Recall for class 1 is 0.89, meaning that 89% of the actual canceled instances were correctly predicted.

The confusion matrix shows that there are some False Positives and False Negatives, but the model is doing a good job of minimizing them.

Additionally, the model is not overfitting, as the metric values for the test and train sets are close together, indicating that the model is generalizing well to unseen data.

```
[77]: # XGBoost Model Building
xgb = XGBClassifier(booster = 'gbtree', learning_rate = 0.1, max_depth = 5,
    ↪n_estimators = 180)
xgb.fit(X_train, y_train)

y_pred_xgb = xgb.predict(X_test)

acc_xgb = accuracy_score(y_test, y_pred_xgb)
conf = confusion_matrix(y_test, y_pred_xgb)
clf_report = classification_report(y_test, y_pred_xgb)

print(f"Accuracy Score of XGBoost Classifier is : {acc_xgb}")
print(f"Confusion Matrix : \n{conf}")
print(f"Classification Report : \n{clf_report}")
```

Accuracy Score of Ada Boost Classifier is : 0.9817409054050276

Confusion Matrix :

```
[[22419   14]
 [  639 12691]]
```

Classification Report :

	precision	recall	f1-score	support
0	0.97	1.00	0.99	22433
1	1.00	0.95	0.97	13330
accuracy			0.98	35763
macro avg	0.99	0.98	0.98	35763
weighted avg	0.98	0.98	0.98	35763

Accuracy Score: 0.9817 Precision for class 0 is 0.97, indicating that 97% of the instances predicted as not canceled were actually not canceled. Recall for class 0 is 1.00, meaning that 100% of the actual not canceled instances were correctly predicted. Precision for class 1 is 1.00, indicating that 100% of the instances predicted as canceled were actually canceled. Recall for class 1 is 0.95, meaning that 95% of the actual canceled instances were correctly predicted.

The model is not overfitting, as the metric values for the test and train sets are close together, indicating that the model is generalizing well to unseen data. Compared to the Random Forest (RF) model, the XGBoost (XGB) model shows a slight improvement in all the performance metrics. Although the improvement is not significant, it is still better at balancing the trade-off between Precision and Recall, leading to a higher F1-score. This indicates that the XGBoost model is a better model for predicting cancellations.


```
[80]: # Compare the model output score
models = pd.DataFrame({
    'Model' : ['Logistic Regression', 'KNN', 'Decision Tree Classifier',
    ↪ 'Random Forest Classifier', 'XgBoost'],
    'Score' : [acc_lr, acc_knn, acc_dtc, acc_rd_clf, acc_xgb]
})

models.sort_values(by = 'Score', ascending = False)
```

```
[80]:
```

	Model	Score
4	XgBoost	0.981741
3	Random Forest Classifier	0.954702
2	Decision Tree Classifier	0.946677
1	KNN	0.890306
0	Logistic Regression	0.809272

```
[81]: px.bar(data_frame = models, x = 'Score', y = 'Model', color = 'Score', template_
    ↪ = 'plotly_dark', title = 'Models Comparison')
```




Among all the tested classifiers, XGBoost had the best performance in predicting hotel booking cancellations. Random forest classifier and Decision Tree Classifier also have good score but lesser compare to XGBoost.

```
[86]: # Feature importance for XGBoost Model classifier
xgb = XGBClassifier()
xgb.fit(X_train, y_train)

feature_importance = pd.Series(xgb.feature_importances_, index=X.columns)
top_features = feature_importance.nlargest(10) # Selecting top 10 features

# Plot feature importance
plt.figure(figsize=(10, 6))
top_features.plot(kind='barh')
plt.title('Feature Importance for XGBoost Model Classifier')
plt.xlabel('Importance')
plt.ylabel('Feature')
plt.show()
```



Feature importance for an XGBoost model classifier is a crucial aspect of understanding the model's behavior and predictive power. By analyzing feature importance, we can identify which variables have the most significant impact on the model's predictions, helping to prioritize resources and interventions accordingly. This insight enables stakeholders to make informed decisions, optimize model performance, and gain deeper insights into the underlying factors driving classification outcomes.